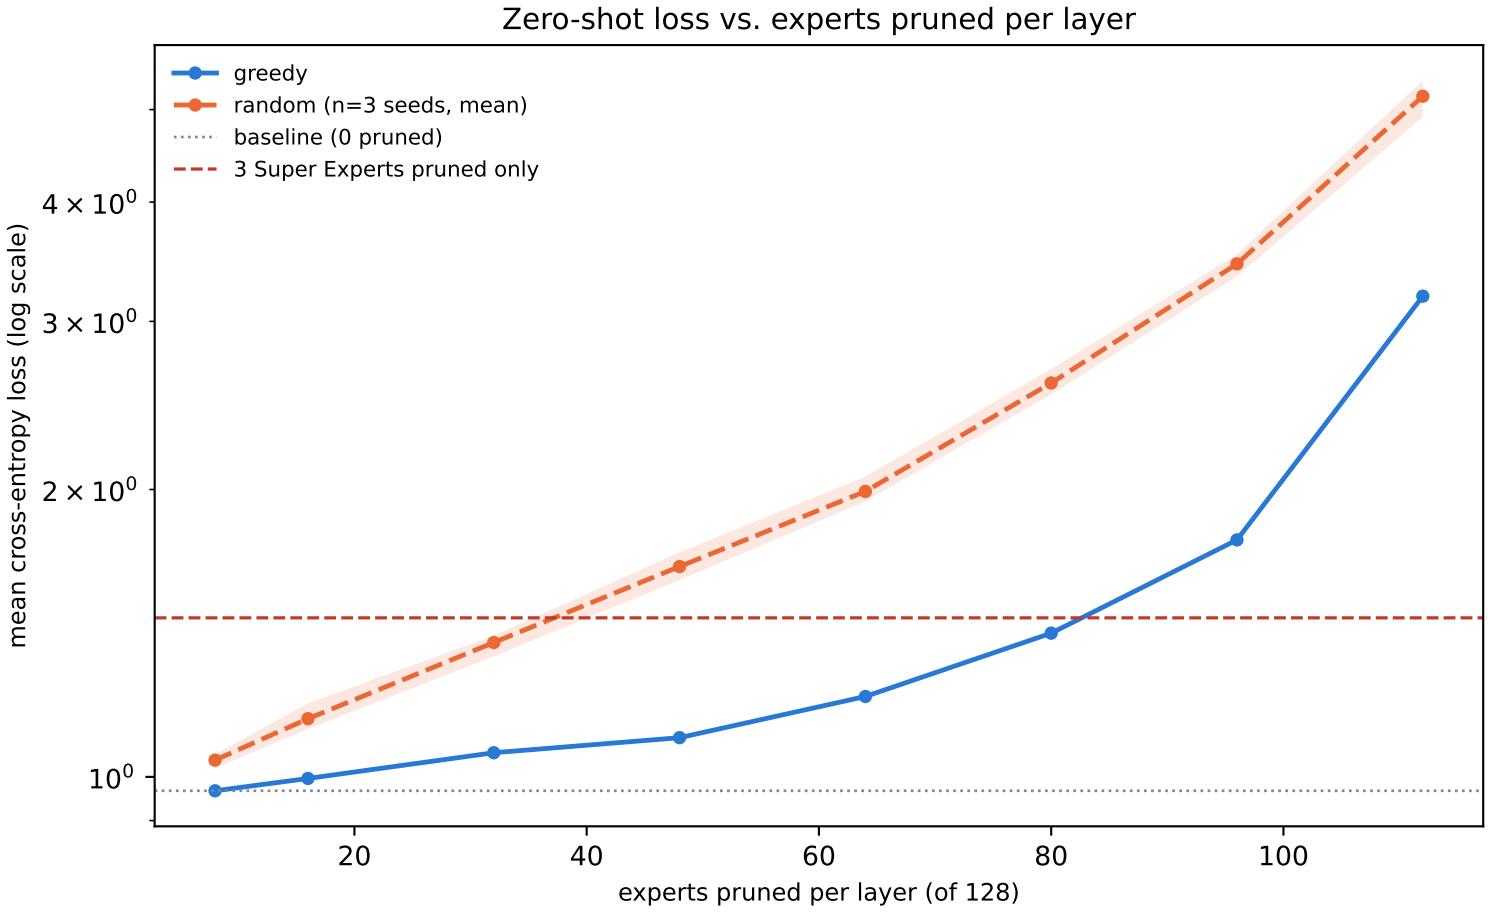
Zero-shot loss vs. experts pruned per layer
greedy
random (n=3 seeds, mean)
baseline (0 pruned)
3 Super Experts pruned only
mean cross-entropy loss (log scale)
4×10^0
3×10^0
2×10^0
10^0
20
40
60
80
100
experts pruned per layer (of 128)