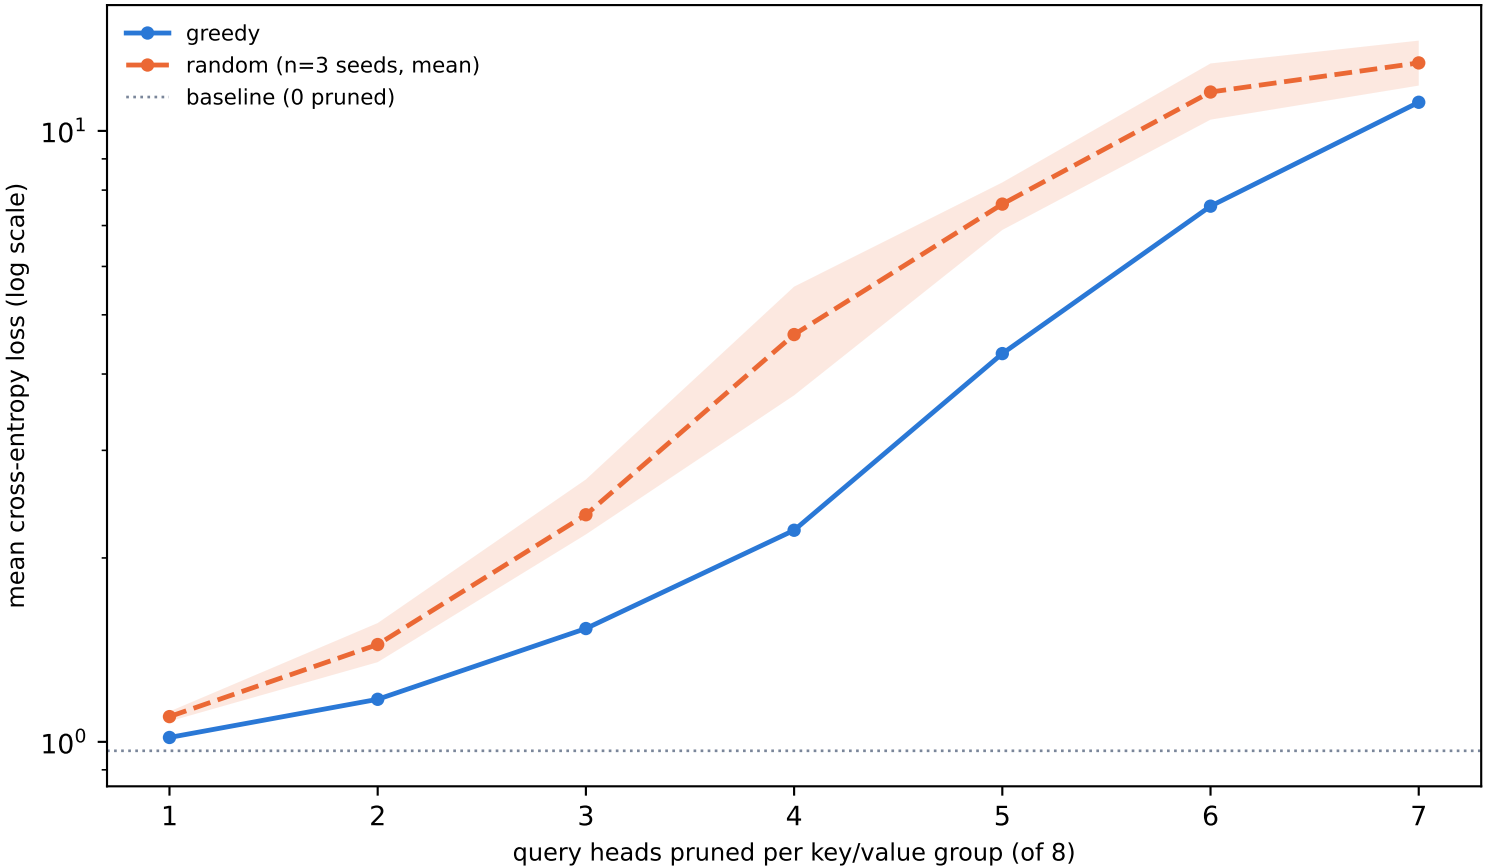
Zero-shot loss vs. query heads pruned per key/value group
greedy
random (n=3 seeds, mean)
baseline (0 pruned)
10^1
10^0
mean cross-entropy loss (log scale)
1
2
3
4
5
6
7
query heads pruned per key/value group (of 8)