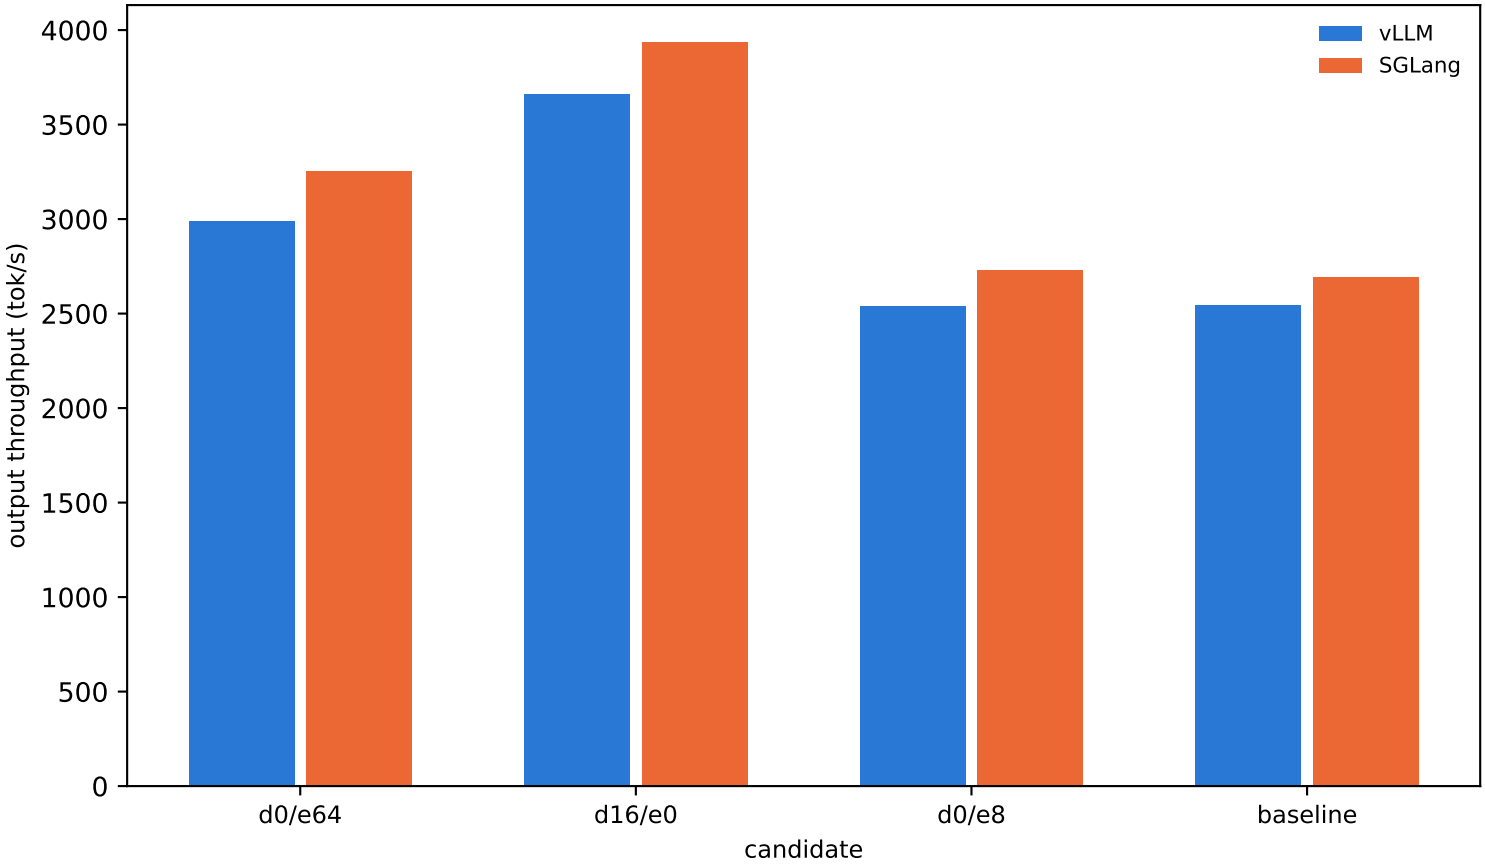
HumanEval-TS serving throughput @ concurrency=16 (512 fixed output tokens)
vLLM
SGLang
output throughput (tok/s)
0
500
1000
1500
2000
2500
3000
3500
4000
d0/e64
d16/e0
d0/e8
baseline
candidate