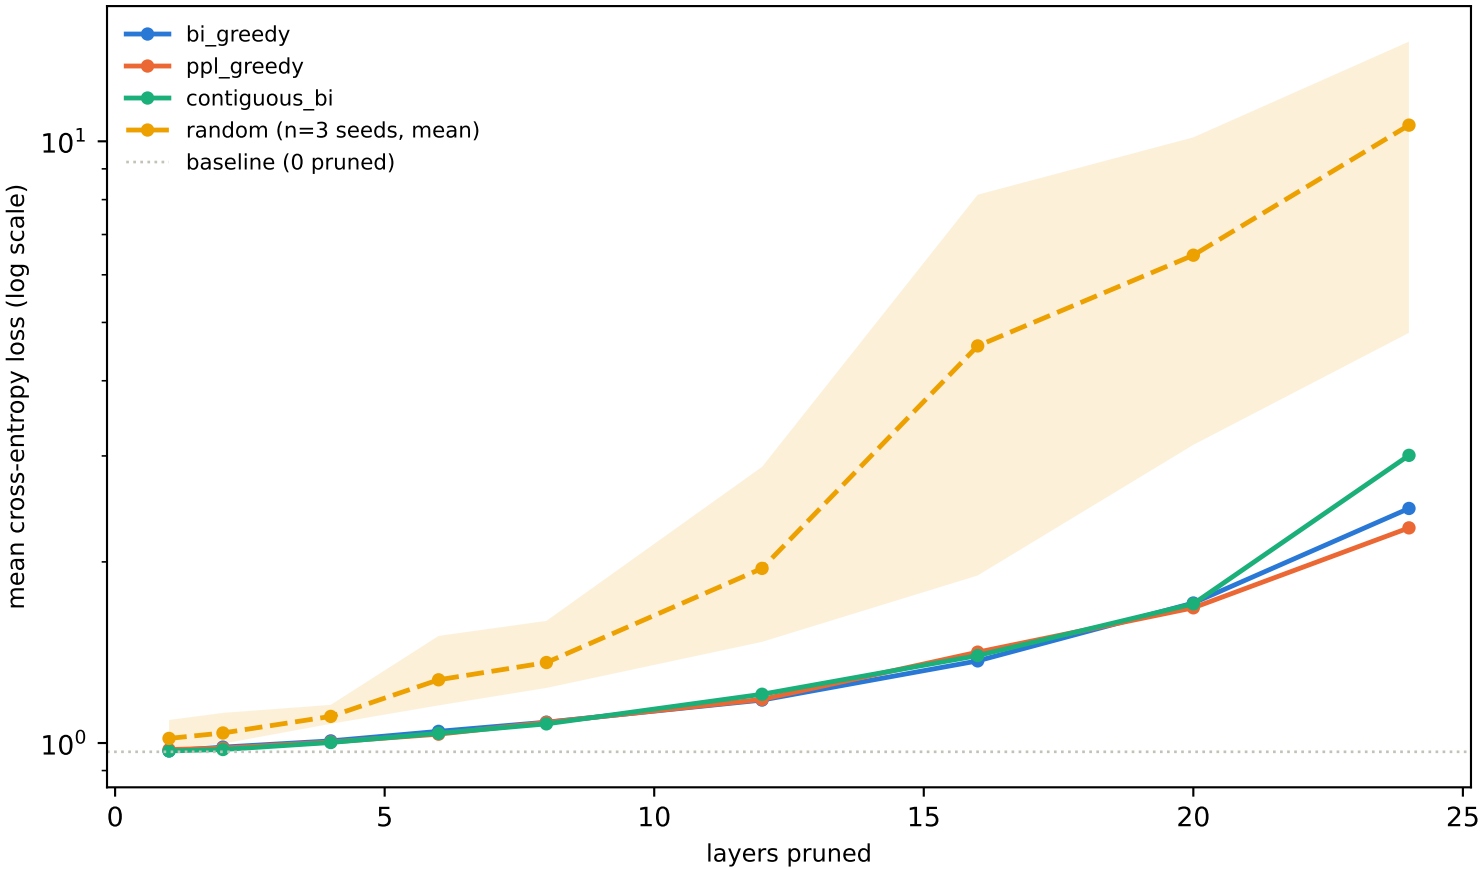
Zero-shot loss vs. layers pruned
bi_greedy
ppl_greedy
contiguous_bi
random (n=3 seeds, mean)
baseline (0 pruned)
mean cross-entropy loss (log scale)
$10^1$
$10^0$
0
5
10
15
20
25
layers pruned